LEGALEDIT: A Dissociation Diagnostic for Legal Meaning Preservation Metrics

Gaurab Baral
University of Cincinnati
baralgb@mail.uc.edu

Abstract

Does a simplified legal clause still say what the original said? The checks in current use cannot establish that it does: requiring an identical pair to score highest and an unrelated pair lowest moves lexical overlap and legal force together, so any monotone function of token overlap satisfies both. Our remedy is a *dissociation*, an item holding surface form fixed while legal force moves. We release LEGALEDIT, 373 minimal perturbations of Quebec statutory French that reverse legal force while preserving 0.93 of the tokens, with a harness that scores metrics, regressors and prompted judges alike. The seven embedding and BERTScore metrics we test spend only 0.022–0.039 of their identical-to-unrelated range on such an edit, against 0.67 for bidirectional NLI, the one family the identical-pair check would disqualify. On FRJUDGE, against a measured human ceiling of $r = 0.597$, a bare length feature outscores every semantic metric and has the lowest margin we measure, while a prompted judge passes both conventional checks perfectly and still spends only 0.370 on a reversed clause, a limit set by the corpus's deduct-per-error rubric rather than by the model.

1 Introduction

Rewriting a legal text for a lay reader has an unusual failure mode. When a news summary drops a qualifier the reader is mildly misinformed; when an insurance clause drops a qualifier the reader may be uninsured. The 2024 decision in *Moffatt v. Air Canada*, which held an airline liable for its chatbot's misstatement of its own bereavement-fare policy, established that the gap between what a document says and what an automated system says it says is legally cognisable. Nor is the scale hypothetical: Hariri and Ho [8] document a U.S. state claiming to have eliminated a third of its code with automated assistance and measure an F1 of 0.67 for such systems, and Magesh et al. [11] find that legal research tools marketed as having eliminated hallucination still produce it on 17 to 33% of queries.

If simplified legal text is produced automatically, something has to check it, and that something will in practice be an automatic metric. What a practitioner needs to know of it is narrower than correlation with human ratings on naturally occurring simplifications: does the score move when the legal effect of the clause moves? Two automated checks are in current use [2, 3]: a sentence paired with itself must receive the maximum score, and a sentence paired with an unrelated one the minimum. Both are necessary and both are easy to state. But moving from an identical pair to an unrelated pair changes lexical overlap and legal meaning *together*, in the same direction, at the same time, so any monotone function of token overlap satisfies both by construction, which is why the embedding metrics in the literature already pass them. The checks test calibration at two endpoints; they do not test what the metric is reading.

The remedy is a *dissociation*: an evaluation item on which the two confounded variables are pulled apart, so that legal force changes while surface form does not. Consider a clause in which *doit* (must) is replaced by *peut* (may). Ninety-odd percent of the tokens are untouched, the sentence remains fluent and statutory in register, and an obligation has become a permission. A metric that measures legal meaning must drop sharply; a metric that measures overlap cannot. We build LEGALEDIT, 373 such perturbations over Quebec statutory French, and release it with a harness taking an arbitrary scoring function, so that untrained metrics, regressors and prompted judges are administered the same test.[1]

LEGALEDIT separates the metric landscape cleanly (Section 5): the similarity metrics collapse,

[1] https://github.com/Nyvora-Vision-Labs/LEGALEDIT

spending 0.022–0.039 of their working range on the edit, while bidirectional NLI entailment spends 0.67. The diagnostic is therefore hard but not impossible, which makes a low score a fixable deficiency rather than a limit of automatic evaluation. Awkwardly for current practice, the family that clears it is the one the identical-pair check would disqualify.

A diagnostic alone does not discipline how corpus-correlation evidence is read, so we assemble three further requirements around it (Section 3): a *measured* human ceiling, comparison under *identical supervision*, and a *trivial-feature control*. We instantiate all four on FRJUDGE [2], 297 Quebec insurance clauses with five expert ratings each and the only annotated corpus for this construct. It supports a human ceiling of $r = 0.597$; a ridge over ten surface features reaches $r = 0.62$, at that ceiling and above every semantic metric we test; and the two orderings disagree at both extremes.

Contributions.

- LEGALEDIT, a 373-item dissociation diagnostic defined by legal effect rather than by linguistic category, with a scoring harness and the margin statistic it is read with.
- An argument that the sanity checks in current use are jointly satisfiable by lexical overlap, and a four-part protocol that repairs this.
- An empirical separation of the field into a similarity family that fails the diagnostic and an entailment family that largely clears it, the first human-ceiling estimate for this construct, and a trivial surface baseline that outscores every semantic metric on the corpus.
- A prompted judge through the same harness, whose diagnostic sensitivity is limited by the corpus's deduct-per-error rubric rather than by the model.

2 Related Work

Evaluating simplification. Metrics borrowed from translation and summarisation transfer poorly to meaning preservation. Sulem et al. [14] show BLEU does not correlate with it once sentence splitting is involved, the operation most characteristic of legal simplification, and embedding metrics [17] capture semantic variability better but score consistently high across *all* alterations, penalising even substantial meaning changes too little. Trained metrics such as MeaningBERT [3] improve correlation on the corpora they are fitted to. FRJUDGE [2] is to our knowledge the only annotated resource for legal meaning preservation in any language, and is the corpus we build on.

Challenge sets and entailment. Our diagnostic follows work showing that correlation on naturally occurring outputs hides brittleness on targeted perturbations. Closest to us, Chen and Eger [4] demonstrate that similarity-based metrics are not robust to meaning-changing edits while NLI-based metrics repair much of the gap, and Mujahid et al. [12] run the complementary experiment, holding meaning fixed and varying form. A usable metric must move on the first kind of edit and stay put on the second, and current metrics reliably do neither. LEGALEDIT adapts the method to a domain where the perturbations are the exact operations that decide coverage disputes, so the inventory is organised by legal effect rather than by linguistic category; to our knowledge no prior challenge set targets legal force. The entailment side of this literature is developed for English adequacy; we construct a French analogue by running a multilingual NLI model [10] in both directions, so that failure of $o \Rightarrow s$ signals hallucination and failure of $s \Rightarrow o$ signals omission, the two dominant categories of the legal error taxonomy.

Whether a benchmark measures what it claims. Reviewing 445 LLM benchmarks, Bean et al. [1] find recurring choices of phenomenon, task and scoring metric that undermine the validity of the claims those benchmarks support, and Pacchiardi et al. [13] give the concrete version: simple n-gram classifiers, possessing none of the capacity under test, score well, so a benchmark can be solved by a route that bypasses the construct entirely. Requirement 3 below is that control transposed to a metric comparison. We follow Deutsch et al. [5] and Graham and Baldwin [7] on statistical practice, since these correlation estimates carry wide intervals and the Williams test is the appropriate one for dependent correlations. And we treat the five ratings per item as a distribution to be modelled rather than noise around a truth, as the perspectivist literature now does by default [15]: legal interpretation is a paradigm case, where disagreement between competent lawyers is the ordinary condition of the field rather than annotator error.

3 What Would Validate a Legal Meaning Metric?

The four requirements below are stated for an arbitrary scorer $m(o, s) \in \mathbb{R}$ over an original clause o and a candidate simplification s. None presuppose an architecture, and all four are cheap enough that we see no reason for a metric paper to omit them.

Requirement 1: a measured ceiling. A correlation with human judgment means nothing until one knows what correlation a human achieves. Where a corpus carries $k \geq 3$ ratings per item this is free: hold out one rater, correlate their ratings against the mean of the remaining $k-1$, and average over the held-out position. Two bounds should be reported. The leave-one-annotator-out correlation answers *is the metric as good as an expert?*; the Spearman–Brown reliability of the k-rater mean, and its square root, answer *is the metric as good as the aggregate label allows?* A metric can sit above the first and below the second without being superhuman, because predicting an average is easier than being one of the raters averaged.

Requirement 2: identical supervision. Comparisons in this literature routinely place an unsupervised cosine similarity and a regressor fine-tuned on the label in the same table, normalise the former by decimal scaling, and report RMSE for both. Decimal scaling fixes the range but neither the location nor the shape of the distribution, so a similarity occupying $[0.7, 1.0]$ records a large RMSE however well it ranks items, and the column measures access to supervision rather than judgment. The repair is to fit a monotone map from each raw score to the human scale on the training split only. We use isotonic regression, which cannot improve a ranking: rank statistics move by at most 0.014 across our comparison set, while Pearson r, which a monotone map need not preserve, moves by up to 0.074.

Requirement 3: a trivial-feature control. Simplification changes surface statistics systematically, and the most systematic is length: LLM simplifications are usually shorter, and the dominant annotated error category in legal simplification is omission. A rubric deducting score once per identified error therefore produces a label partly predictable from how many words disappeared, with no representation of legal force involved. We require a supervised baseline over surface features alone (word counts, their difference, token and numeral overlap), fitted under Requirement 2: on the precedent of Pacchiardi et al. [13], a comparison without one has not established that its scores require the construct.

Requirement 4: dissociation. The three requirements above discipline the corpus-correlation evidence but cannot establish what a metric is reading, because on naturally occurring pairs surface form and legal force vary together. This is the confound that makes the identical and unrelated checks uninformative, and it needs the same remedy: a set on which the two are pulled apart by construction. Let s' be a perturbation of s with token Jaccard $J(s, s') \approx 1$ and legal force $L(s') \neq L(s)$, and let u be an unrelated sentence of the same register. We report the *margin fraction*

$$\text{margin} \;=\; \frac{\overline{m(s,s)} - \overline{m(s,s')}}{\overline{m(s,s)} - \overline{m(s,u)}},$$

the share of the metric's own identical-to-unrelated range that it spends on a legally decisive edit, which charges each metric against the range it actually has rather than a nominal $[0, 1]$. A metric that treats the flipped sentence as it treats an unrelated one scores 1; a pure overlap function scores near 0. The margin is not the rate at which a metric ranks (s, s) above (s, s'), and the rate is the weaker quantity, since a metric can rank correctly on every item while moving by a thousandth of its range.

4 LegalEdit: Constructing the Dissociation

Source text and edit inventory. We draw source sentences from two Quebec statutes held out from any corpus used elsewhere in the paper, the Automobile Insurance Act and the Highway Safety Code, and apply a single minimal edit that changes legal force while leaving the sentence otherwise intact. The inventory is organised by legal effect rather than by linguistic category, which is where it departs from general-domain perturbation suites. Six families make up *tier 1*, in which the edit unambiguously changes who is bound, what is covered or how much: modality (*doit* $\leftrightarrow$ *peut*, obligation becomes permission), party (*l'assuré* $\leftrightarrow$ *l'assureur*, the duty changes hands), polarity (dropping *ne . . . pas*, *exclut* $\leftrightarrow$ *inclut*), quantum (the first amount multiplied by ten), scope (*tous les* $\rightarrow$ *certains*, *sauf si* $\rightarrow$ *si*) and temporal (*avant* $\leftrightarrow$ *après*). A *tier*

2 set of connective and class shifts (*automobile* ↔ *véhicule*, *et* ↔ *ou*) is reported separately. The supplementary material gives one example per rule.

Keeping the edit invisible to a fluency detector. A diagnostic of this kind fails if a metric can find the edited item by noticing that it reads badly. Perturbations introducing a grammatical artefact are therefore filtered by a conservative French well-formedness check covering failed elision (*le automobile*), missing contraction (*de le*), doubled negation and determiner–noun gender disagreement; it flags none of the 400 unmodified source sentences, so it removes only genuine artefacts. The result is 373 pairs (310 tier 1) with mean token Jaccard 0.93 against their source, none below 0.75.

Administering it, and to what. For each source sentence s with perturbation s' we score (s, s), (s, s') and (s, u) for an unrelated u from the same statutes, min–max normalise per metric so the identical and unrelated probes anchor the range, and compute the margin from the three means. We administer it to French-capable metrics throughout, since scoring English-only metrics on French text confounds language coverage with metric quality. The set covers three families, named in full in Table 1: BERTScore [17] over CamemBERTv2 and FlauBERT, in F1 and recall-only variants since recall targets omission; four sentence-embedding cosines; and bidirectional NLI entailment with mDeBERTa-XNLI [10]. To these we add token Jaccard and the trivial surface features of Requirement 3, so that the overlap end of the spectrum is represented explicitly.

5 Diagnostic Results: Two Families, Not One

Table 1 reports the margin over 13 scorers, plotted per metric in the supplementary material. The metrics divide cleanly into two families.

Every *similarity*-based metric fails, and fails almost completely. The seven embedding and BERTScore variants spend between 0.022 and 0.039 of their working range on an edit that reverses the law. BERTScore over CamemBERTv2 places a sentence saying the opposite of the original at 0.982, on a scale where the original itself sits at 1.000 and an unrelated statute sits at 0.193; token Jaccard, at 0.079, is in the same regime, which is the point: these metrics are behaving as overlap functions, exactly as Section 3 says they must. Their *discrimination* rates are uninformative for the same reason, since the seven rank the identical pair above the edited one 100% of the time but by a margin that carries no decision.

Bidirectional NLI does not fail. It spends 0.67 of its range on the edit, ranks correctly on 93.3% of items, and places a legally flipped sentence (0.265) nearer the unrelated pair (0.010) than the identical one (0.784); the two directions individually reach 0.424 and 0.473, and their minimum, which an omission-*or*-hallucination reading of the construct requires, produces the 0.67. That entailment survives meaning-changing edits better than similarity is not itself new [4]; what is new is an inventory defined by legal effect and the observation below.

The existing checks would have excluded it. Note where NLI's identical score sits: 0.784, not 1.000, and not as a normalisation artefact. In raw terms mDeBERTa assigns a mean entailment of 0.780 to a sentence paired with *itself*, clearing the 99% bar the identical-pair convention requires on only 1.3% of items, against 100% for every BERTScore and embedding metric. The metric that best tracks legal meaning is the one metric convention would *disqualify*. That check does not merely fail to identify a legal metric, it selects against the family with the property being sought. Saturation on identical input is easy to build, since every overlap function has it by construction, and is evidence about calibration at one endpoint rather than about whether the metric reads the law. It should be reported, and it should not be a filter.

No rule is reliably detected. The best-detected operation, dropping *ne . . . pas*, averages only 0.333, while swapping the bound party draws 0.035 against 0.155 across all rules: reversing which party a clause binds is close to the least detectable thing one can do to it. The supplementary material gives the per-rule breakdown.

6 Case Study: The Protocol on FrJudge

The diagnostic answers one of the four requirements. This section runs the other three on FrJudge [2], 297 pairs drawn from two Quebec insurance forms, each simplified by `gpt-4-turbo-2024-04-09` with a zero-shot prompt and rated by five law students on legal meaning (1–10), simplicity, and an 18-class legal characterization. It then asks whether the two kinds of evidence agree. The public `FrJUDGE.jsonl` drops annotator identity, which makes Requirement

Metric	identical	LEGALEDIT	unrelated	discrim.%	margin
NLI-bidirectional (min)	0.784	0.265	0.010	93.3	0.670
NLI-bwd (no omission)	0.784	0.432	0.042	77.2	0.473
NLI-fwd (no hallucination)	0.784	0.469	0.043	77.7	0.424
LLM-judge (deepseek-chat, k=5)	1.000	0.632	0.006	86.3	0.370
surface: token Jaccard	1.000	0.928	0.089	99.2	0.079
LaBSE	1.000	0.972	0.292	100.0	0.039
BERTScore-FlauBERT	1.000	0.979	0.393	100.0	0.035
mE5-base	1.000	0.975	0.266	100.0	0.034
Para-mUSE-mpnet	1.000	0.981	0.312	100.0	0.028
Sentence-CamemBERT	1.000	0.983	0.316	100.0	0.024
BERTScore-CamemBERTv2-Recall	1.000	0.983	0.240	100.0	0.023
BERTScore-CamemBERTv2	1.000	0.982	0.193	100.0	0.022
surface: $-\|\,\|s\| - \|o\|\,\|$	1.000	0.997	0.718	11.3	0.011

Table 1: LEGALEDIT. Scores min–max normalised per metric. *discrim.* is the share of items ranked correctly (identical $>$ edited); *margin* is the fraction of the identical-to-unrelated range spent on the legal edit. The similarity metrics rank 100% of items correctly while moving almost nothing, which is why the margin is the quantity to read.

1 impossible for a downstream user, so we rebuild the corpus from the raw Prodigy export, recovering the five annotators (pseudonymised A–E) and the source document of each clause, and reproduce the published mean pairwise agreement of 25.96% exactly.

6.1 Disagreement, and the human ceiling

Krippendorff's α for legal meaning is 0.104 under the *nominal* coefficient, conventionally the one reported for this corpus, and 0.325 under the *ordinal* one. Nominal α counts a one-point disagreement on an ordered ten-point scale as fully as a nine-point one, so the ordinal coefficient is the one appropriate to the design [9]: the corpus is more reliable than a nominal reading suggests, though still only fairly. The disagreement that remains is structural rather than random. The five annotators separate into two stable populations, B, C and D averaging 7.61 against 4.34 for A and E; ratings are strongly bimodal, the median across-annotator range on a pair is 7 points, and 63% of pairs span at least six. Averaging produces a target with a standard deviation of only 2.40, precisely the condition under which a regressor that predicts the middle scores well.

Requirement 1 is directly executable here. Correlating each annotator's ratings against the mean of the other four gives a mean leave-one-annotator-out correlation of 0.597 (RMSE 3.27), with a mean pairwise correlation of 0.462. A model exceeding 0.597 is not thereby superhuman: predicting the *mean* of five raters is a much easier task than being one of them, and rewards predicting the middle of a bimodal distribution. The upper bound for predicting that mean is its own reliability, 0.811 by Spearman–Brown, so a metric can in principle approach $\sqrt{0.811} = 0.90$. A correlation between 0.597 and 0.90 is evidence about the aggregate label and about nothing at all with respect to matching an expert.

6.2 Metrics under identical supervision

Requirement 2 asks for a comparison in which every candidate sees the same supervision, and Requirement 3 for a trivial control inside it. We take the comparison set of Section 4 and add a supervised ridge regression over ten surface features, mapping every metric's raw score to the human scale by an isotonic regression fitted on the training split alone.

Two results follow (the full table is in the supplementary material). First, calibration collapses most of the RMSE gap that uncalibrated comparisons record: the metrics were not as far from human judgment as such tables suggest, merely on a different scale. Second, a ridge over word counts and token overlap reaches $r = 0.62$, at the human ceiling and above every semantic metric we test, the best of which reaches only $r = 0.482$. Only two scorers beat the ridge, and neither is semantic: the prompted judge of Section 6.3 and the bare length difference at $r = 0.641$, the trivial feature the ridge was built around. Whatever a fine-tuned legal metric adds beyond that, a comparison without this row does not measure it. This is evidence not that the corpus is trivial but about what the *annotation* rewards: score is reduced once per identified error, the most frequent error type is omission, and omission correlates strongly with becoming shorter. A metric can therefore score well by detecting deletion, with no representation of legal force, which is the hypothesis Requirement 4 exists to test. The test split holds 89 items, at which differences finer than roughly a quarter of a correlation point are not resolvable, ours included; we report bootstrap intervals throughout [5], and under the Williams test the length feature's advantage is significant against each of the ten unsupervised metrics.

6.3 A prompted judge

A prompted model given the annotators' own rubric [6] is the natural competitor to a fine-tuned 112M-parameter encoder. We close that gap with `deepseek-chat` under the FRJUDGE scale as the corpus describes it: a ten-point judgment starting at 10 and deducting one point per identified legal error, each pair scored in its own request so the judge never sees the probe condition, $K = 5$ times with the mean reported. Calibrated under Requirement 2 it reaches $r = 0.755_{\pm 0.054}$, above the length feature, the surface regression and our ceiling of 0.597, although an independent redraw moved it to 0.726, so it is better quoted as a range. On LEGALEDIT it scores the identical pair at 10.00 and the unrelated pair at 1.06, a flawless result on precisely the two checks Section 3 argues are jointly satisfiable by lexical overlap, and then awards a clause whose legal force has been reversed 6.69 out of ten, a margin of 0.370.

The rubric, not the model, is what limits it. Holding the model fixed and varying only the wording, the margin moves from 0.421 under the rubric-faithful prompt to 0.687 under a bare instruction that strips the four-type taxonomy, a range of 0.267 across 3 wordings of one question, against 0.002 for an independent redraw of the same prompt. The mechanism is visible in the rubric: score is reduced *once per identified error*, so a clause reversal is one incoherence and the scheme caps the penalty near the top of the scale. A judge applying the corpus's own rubric faithfully is thereby prevented from expressing that an edit destroys the legal meaning. Neither capability nor combination closes the gap: a reasoning model under the identical rubric reaches margin 0.320 against 0.425, and of six rules for merging the judge with bidirectional NLI none beats both parents on both axes. One behaviour survives every configuration: substituting *peut* for *doit* scores 4.72 while the reverse scores 8.02, a gap of 3.30, though both edits change the law by the same construction. A scorer that treats a stricter rewrite as more faithful is tracking legal

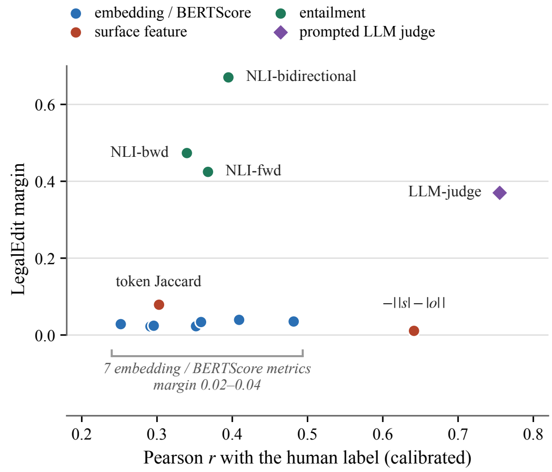

Figure 1: The two orderings come apart. Horizontally, correlation with the human label after calibration, the quantity a conventional comparison selects on; vertically, the share of working range spent on an edit that reverses the law. Selecting on the horizontal axis picks the prompted judge and then the length feature, the worst scorer of all on the vertical axis.

severity rather than legal *equivalence*. The supplementary material carries the per-configuration table and these controls in full.

6.4 The two orderings come apart

Over the 13 scorers that carry both a calibrated r on FRJUDGE and a LEGALEDIT margin, the rank correlation between the two is $\rho = 0.24$ ($p = 0.43$), with a 95% bootstrap interval of $[-0.39, 0.75]$. That interval is an absence of resolution rather than a demonstrated null: we are not entitled to say that corpus correlation fails to predict legal sensitivity, only that 13 metrics cannot establish that it does. Restricting to the 10 untrained neural baselines moves the estimate to $\rho = 0.58$ (95% CI $[-0.08, 0.90]$), pointing the other way, but that subset excludes the length feature and the prompted judge, precisely the two scorers at which the orderings part.

The extremes are the durable part of this result and do not depend on a rank correlation over thirteen points. The trivial length feature outscores every semantic metric ($r = 0.641$) and has the lowest margin of anything we test (0.011), while bidirectional NLI sits mid-table on correlation and is the only family that reacts to the edit at all: a metric selected on corpus correlation alone would pick the wrong one. The prompted judge makes the same point from the other end, strongest on the corpus and fourth of 13 on the diagnostic. Figure 1 plots the two axes against each other, and the emptiness of its upper-right quadrant is the finding.

7 Discussion

Three checks, none of them a filter. A metric claiming to measure legal meaning should report three checks, not two: identical pairs (does it saturate?), unrelated pairs (does it bottom out?), and LEGALEDIT (does it move when, and only when, the law moves?). The first two are jointly satisfiable by lexical overlap; the third cannot be passed by overlap at all. That LEGALEDIT is passable matters as much as that it is hard: a low margin identifies a fixable deficiency, not a limit of automatic evaluation. None of the three should be an admission criterion, since here the metric that scores highest on the diagnostic fails the identical-pair check most badly. Saturation, bottoming out and legal sensitivity are independent properties, and a scorer can be repaired for the first two by calibration and cannot be repaired for the third at all. Rank on the property that cannot be fixed.

Entailment is a starting point, not a destination. The construct FRJUDGE annotates is defined by an error taxonomy whose two dominant categories are omission and hallucination, precisely the two directions of entailment failure, and the only scorer that reacts to a legally decisive edit is built from bidirectional entailment. That does not make an off-the-shelf NLI model an adequate legal metric: its correlation is middling ($r = 0.395$), its margin leaves a third of the range unspent, and a simplification that glosses a defined term adds content the clause does not contain, which the $o \Rightarrow s$ direction reads as hallucination. The narrower conclusion is that the next legal-meaning metric should be built on a repaired entailment objective and evaluated on the diagnostic, not on a similarity objective selected by corpus correlation. Nor should the disagreement behind the label be averaged away: two of five law students rate legal preservation three points lower than the other three, and a metric that reports "6.2, experts ranged 2–9" tells a practitioner that this clause needs a human, which "6.2" does not.

8 Conclusion

Legal meaning preservation is the right construct to evaluate, but the evidence a metric paper in this area conventionally offers cannot establish that a metric measures legal meaning rather than lexical overlap and length. The missing ingredient is a dissociation, and it is cheap: hold surface form fixed and move the law. On that test every similarity-based metric we examined scores a sentence saying the opposite of the original as though nothing had changed. The test is not unpassable: bidirectional entailment, a family the standard comparison leaves out and the standard sanity check would have disqualified, clears most of it. It also warns about metric selection, because a bare length feature outscores every semantic metric on the corpus and ranks last on the law.

Limitations

Our LEGALEDIT labels are definitional rather than annotated: we assert that changing *doit* to *peut* changes legal meaning, rather than having lawyers rate each perturbation. We think this is defensible for tier-1 edits and report tier-2 shifts separately for exactly this reason, but a human-validated version would be stronger and a small expert study is the obvious next step. The perturbations are also single-edit and template-generated, so a system could in principle be tuned to them without acquiring general legal sensitivity, which is why we treat LEGALEDIT as a necessary condition and not a benchmark to be optimised.

That caveat bears on our most encouraging result. Our perturbations are largely negation, modality, quantifier and antonym contrasts, close to the inventory natural language inference training data is built around, so bidirectional NLI's margin of 0.67 may reflect familiarity with the *form* of our edits rather than sensitivity to their legal consequence. Yang et al. [16] put the general version: any counterfactual edit is a compound treatment bundling the variable of interest with incidental surface variation, and belongs to the target factor only once compared against a meaning-preserving paraphrase baseline. Our margin is normalised against each metric's own range and the identical probe is a null-edit anchor, but we run no paraphrase-control arm. Little turns on this for the similarity metrics, whose margins leave nothing for a control to explain away; the objection bites on bidirectional NLI, and the next version of the diagnostic should carry legally inert paraphrases and defined-term substitutions alongside the decisive edits.

Three further limits. The diagnostic and the corpus comparison run on different text, insurance forms against two statutes, so they answer different questions and should not be read as one ranking. Our judge results rest on one model from one vendor under a rubric reconstructed from the published description rather than the annotators' instructions, which given how far 0.267 of margin turns on wording is a load-bearing assumption. And we evaluate isolated clauses from a single generator, with the ceiling computed on all 297 items while model correlations run on 89-item splits, so "model > ceiling" should be read with the sampling uncertainty of both in view.

Ethical Considerations

FRJUDGE's five annotators are pseudonymised as A–E throughout, our released code pseudonymises annotator identifiers on load, and we ship the pseudonymisation script separately for anyone redistributing a reconstruction. Our results also strengthen a caution Beauchemin et al. [2] raise: a metric that cannot distinguish *doit* from *peut* must not gate whether a simplified clause is safe to publish, and the appropriate deployment is triage, the ranking of clauses for human review, rather than certification. Nothing here constitutes legal advice. The perturbed sentences are deliberate misstatements of Quebec law, are labelled as such in the distribution, and should not be quoted as statements of the law.

References

[1]. Andrew M. Bean, Ryan Othniel Kearns, Angelika Romanou, et al. 2025. Measuring what matters: Construct validity in large language model benchmarks. In *Advances in Neural Information Processing Systems 39 (NeurIPS) Datasets and Benchmarks Track*.

[2]. David Beauchemin, Michelle Albert-Rochette, Richard Khoury, and Pierre-Luc Déziel. 2025. JUDGEBERT: Assessing legal meaning preservation between sentences. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 92–118.

[3]. David Beauchemin, Horacio Saggion, and Richard Khoury. 2023. MeaningBERT: assessing meaning preservation between sentences. *Frontiers in Artificial Intelligence*, 6:1223924.

[4]. Yanran Chen and Steffen Eger. 2023. MENLI: Robust evaluation metrics from natural language inference. *TACL*, 11:804–825.

[5]. Daniel Deutsch, Rotem Dror, and Dan Roth. 2021. A statistical analysis of summarization evaluation metrics using resampling methods. *TACL*, 9:1132–1146.

[6]. Joseph Enguehard, Morgane Van Ermengem, Kate Atkinson, et al. 2025. LeMAJ (legal LLM-as-a-judge): Bridging legal reasoning and LLM evaluation. In *Proceedings of the Natural Legal Language Processing Workshop 2025*, pages 318–337.

[7]. Yvette Graham and Timothy Baldwin. 2014. Testing for significance of increased correlation with human judgment. In *EMNLP*.

[8]. Emaan Hariri and Daniel E. Ho. 2025. AI for statutory simplification: A comprehensive state legal corpus and labor benchmark. In *Proceedings of the Twentieth International Conference on Artificial Intelligence and Law (ICAIL)*.

[9]. Klaus Krippendorff. 2004. *Content Analysis: An Introduction to Its Methodology*. Sage.

[10]. Moritz Laurer, Wouter van Atteveldt, Andreu Casas, and Kasper Welbers. 2024. Less annotating, more classifying: Addressing the data scarcity issue of supervised machine learning with deep transfer learning and BERT-NLI. *Political Analysis*, 32(1):84–100.

[11]. Varun Magesh, Faiz Surani, Matthew Dahl, Mirac Suzgun, Christopher D. Manning, and Daniel E. Ho. 2025. Hallucination-free? Assessing the reliability of leading AI legal research tools. *Journal of Empirical Legal Studies*, 22(2):216–242.

[12]. Zain Muhammad Mujahid, Dustin Wright, and Isabelle Augenstein. 2026. Stress testing factual consistency metrics for long-document summarization. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics*.

[13]. Lorenzo Pacchiardi, Marko Tesic, Lucy G. Cheke, and José Hernández-Orallo. 2024. Leaving the barn door open for Clever Hans: Simple features predict LLM benchmark answers. *arXiv preprint arXiv:2410.11672*.

[14]. Elior Sulem, Omri Abend, and Ari Rappoport. 2018. BLEU is not suitable for the evaluation of text simplification. In *EMNLP*, pages 738–744.

[15]. Yinuo Xu and David Jurgens. 2026. Beyond consensus: Perspectivist modeling and evaluation of annotator disagreement in NLP. *arXiv preprint arXiv:2601.09065*.

[16]. Zihao Yang, Mosh Levy, Yoav Goldberg, and Byron C. Wallace. 2026. Compared to what? Baselines and metrics for counterfactual prompting. In *Conference on Language Modeling (COLM)*.

[17]. Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q Weinberger, and Yoav Artzi. 2020. BERTScore: Evaluating text generation with BERT. In *ICLR*.